


Sample efficiency: gradient descent through the solver is last at a small budget and first at a large one

(a) every arm, same start, same budget



(b) the same curves, zoomed to where the crossings are

