




(a) one spike in eleven moves -- 8.9% of the raster; one beat per class



(b) the start device answers F for every beat



(c) balanced, not the majority class



Same weights, same 16 MIT-BIH beats, same seed. The only difference is the five numbers the FeFET hands up -- th_th 5.622 $\rightarrow$ 4.611 as the memory window widens.

The population rate barely moves (0.4344 $\rightarrow$ 0.4572) and the per-neuron rate profile is unchanged ($r = 0.9999$); what moves is WHICH spikes, and that is enough to unstick the readout.

Loss 1.3996 $\rightarrow$ 1.0177, accuracy 0.250 $\rightarrow$ 0.688, nothing trained in between. The split is INTRA-PATIENT -- record identity is stripped by the thesis' preprocessing, so an

inter-patient AAMI DS1/DS2 split cannot be built from these files at all -- and the inner training loop is deliberately cheap, so these accuracies are NOT comparable to the thesis' fully trained 0.793.